


IT5002 Computer Systems and Applications

Tutorial 10

1. [Open File Table] Below is an illustration of how an OS organizes file data at runtime:



Discuss how this organization helps OS to handle the following scenarios. Your answer should refer to the relevant structure(s) if possible.

- a. Process A tries to open a file that is currently being written by Process B.
- b. Process A tries to use a bogus file descriptor in a file-related system call.
- c. Process A can never "accidentally" access files opened by Process B.
- d. Process A and Process B reads from the same file. However, their reading should not affect each other.
- e. Redirect Process A's standard input / output, e.g. "**a.out < test.in > test.out**". (Hint: *nix processes has 3 default file descriptors: 0 = stdin, 1 = stdout, 2 = stderr)

2. [Putting it together] This question is based on a "simple" file system built from the various components discussed in lecture 12.

Partition Information (Free Space Information)

- Bitmap is used, with a total of 32 file data blocks (1 = Free, 0 = Occupied):

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	0	0	0	0	1	0	0	1	0	0	0	0	1	0	1
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
1	0	1	0	1	0	1	0	1	1	1	0	0	1	0	0

Directory Structure + File Information:

- Directory structures are stored in 4 "directory" blocks. Directory entries (both files and subdirectories) of a directory are stored in a single directory block.
- Directory entry:
 - o **For File:** Indicates the first and last data block number.
 - o **For Subdirectory:** Indicates the directory structure block number that contains the subdirectory's directory entries.
 - o The "/" root directory has the directory block number 0.

0			1			2			3		
y	Dir	3	g	File	0 31	k	File	6 6	i	File	1 3
f	File	12 2	z	Dir	2				h	File	27 28
x	Dir	1									

File Data:

- Linked list allocation is used. The first value in the data block is the "next" block pointer, with "-1" to indicate the end of data block.
- Each data block is 1 KB. For simplicity, we show only a couple of letters/numbers in each block.

0	1	2	3	4	5	6	7
11 AL	9 TH	-1 S!	-1 ND	23 GS	-1 SO	-1 :)	10 TE
8	9	10	11	12	13	14	15
31 RE	3 EE	28 M:	31 OH	19 SE	13 AH	4 IN	17 NO
16	17	18	19	20	21	22	23
30 YE	2 OU	1 ON	17 RI	26 EV	14 AT	21 DA	7 YS
24	25	26	27	28	29	30	31
-1 HO	18 ME	0 AL	30 OP	-1 -(	5 LO	21 ER	-1 A!

- a. (Basic Info) Give:
 - The current free capacity of the disk.
 - The current user view of the directory structure.
 - b. (File Paths) Walkthrough the file path checking for:
 - `"/y/i"`
 - `"/x/z/i"`
 - c. (File access) Access the entire content for the following files:
 - `"/x/z/k"`
 - `"/y/h"`
 - d. (Create file) Add a new file `"/y/n"` with 5 blocks of content. You can assume we always use the free block with the smallest block number. Indicate **all changes required to add the file**.
3. Most OSes perform some higher-level I/O scheduling on top of just trying to minimize hard disk seeking time. For example, one common hard disk I/O scheduling algorithm is described below:
- i. User processes submit file operation requests in the form of
operation(starting hard disk sector, number of bytes)
 - ii. The OS sorts the requests by hard disk sector.
 - iii. The OS merges requests that are nearby into a larger request, e.g. several requests asking for tens of bytes from nearby sectors merged into a request that reads several nearby sectors.
 - iv. The OS then issues the processed requests when the hard disk is ready.
- a. How should we decide whether to merge two user requests? Suggest two simple criteria.
 - b. Give one advantage of the algorithm as described.
 - c. Give one disadvantage of the algorithm and suggest one way to mitigate the issue.
 - d. Strangely enough, the OS tends to **intentionally delay** serving user disk I/O requests. Give one reason why this is actually beneficial using the algorithm in this question for illustration.

- e. In modern hard disks, algorithms to minimise disk head seek time (e.g. SCAN variants, FCFS etc.) are **built into the** hardware controller. i.e. when multiple requests are received by the hard disk hardware controller, the requests will be reordered to minimise seek time. Briefly explain how the high-level I/O scheduling algorithm described in this question may conflict with the hard disk's built-in scheduling algorithm.